

How Action Collapse Creates Illusory Safety in VLA Collision Avoidance Evaluation

Anonymous Author(s)

Affiliation

Address

email

Abstract: Vision-Language-Action (VLA) models report high Collision Avoidance Rates (CAR) on safety benchmarks, yet we show these metrics are contaminated by action collapse, where near-zero motor output produces zero collisions through inaction rather than avoidance. On SafeLIBERO, an SDF-augmented Pi0 achieves 91.5% CAR while 99.6% of its safe episodes exhibit end-effector (EE) displacement below 0.1 mm. The robot is action-collapsed, not safe. We introduce displacement-filtered CAR (dfCAR), which excludes sub-threshold episodes. The bimodal displacement distribution makes dfCAR robust across a [2, 200] mm threshold range. We further show that end-effector-only Signed Distance Field (SDF) supervision fails to deliver collision avoidance. At inference time, ground-truth oracle refinement reduces CAR below unrefined baselines because the SDF gradient drives unmonitored arm links into obstacles. During training, collision-free demonstrations alone account for nearly all measured avoidance gains (~26 percentage points). SDF supervision shows no statistically significant benefit at the tested scale $\lambda \in \{0.005, 0.01, 0.1\}$ ($p > 0.05$ for $\lambda \in \{0.01, 0.1\}$; $\lambda = 0.1$: $57.00\% \pm 2.18$, 3 seeds). Displacement filtering is necessary for meaningful VLA safety evaluation. EE-only geometric supervision does not significantly improve collision avoidance across the tested loss-weight range, suggesting that collision avoidance primarily arises from demonstration data rather than EE-only geometric supervision.

Keywords: VLA safety, collision avoidance, action collapse, signed distance fields

1 Introduction

Collision avoidance is a prerequisite for deploying Vision-Language-Action (VLA) models in physical environments [1]. The SafeLIBERO benchmark [2] provides the first systematic protocol for measuring collision avoidance in VLA manipulation. Recent methods report high Collision Avoidance Rate (CAR) through constrained optimization [3], diffusion-process collision costs [4], and collision-free demonstrations [5]. Our SDF-augmented Pi0 [6] reaches 91.5% CAR across 800 episodes. The problem appears close to solved.

It is not (Figure 1). Of the 732 collision-free episodes, 99.6% exhibit displacement below 0.1 mm; TSR is 0% across all 800 episodes. The robot is action-collapsed, not safe. The same pattern holds across all Pi0 configurations (Section 4). We term this *action-collapsed safety* (colloquially, “fake safety”): *action collapse* (near-zero actions producing zero collisions) renders reported CAR illusory, and standard evaluation cannot distinguish inaction from deliberate avoidance [7].

We present a systematic investigation of this problem across the Pi0 and Pi0.5 [6] model families on SafeLIBERO, spanning 100+ experiments and 6000+ evaluation episodes. Three findings emerge.




Figure 1: The action-collapsed safety problem in VLA collision avoidance evaluation. Standard CAR classifies an action-collapsed robot as safe because it produces zero collisions (left). dfCAR filters out action-collapsed episodes where displacement falls below τ_d ; Pi0+A1’s 91.5% safety score drops to $\approx 0\%$ (center). Pi0.5 exhibits genuine avoidance with active motion above 288 mm in all safe episodes, and dfCAR equals raw CAR (right).

Action collapse is pervasive. We introduce displacement-filtered CAR (dfCAR), which excludes episodes below a displacement threshold τ_d . dfCAR collapses Pi0+A1 from 91.5% to $\approx 0\%$ (3/732 safe episodes survive). The bimodal displacement distribution makes τ_d robust across five orders of magnitude. Pi0.5 validates the metric: TSR=71.5%, CAR=28.5%, with every safe episode above 288 mm displacement and dfCAR equal to raw CAR (Section 4).

EE-only SDF refinement is structurally counterproductive. Three ground-truth oracle variants all achieve *lower* CAR than unrefined baselines (73.1–76.9% vs. 82–83%). The SDF gradient repels the EE but rotates arm links into obstacle volumes (Section 5).

Collision-free demonstrations are sufficient for collision avoidance. LoRA fine-tuning on Pi0.5 with collision-free demonstrations raises CAR by ~ 26 points with zero action-collapsed episodes. Adding SDF supervision does not yield statistically significant improvement across the tested range $\lambda \in [0.005, 0.1]$ ($p > 0.05$ for $\lambda \in \{0.01, 0.1\}$; $\lambda = 0.1$: $p = 0.131$, 3 seeds). The per-task avoidance ranking is preserved across all conditions, indicating that the demonstration distribution primarily accounts for observed avoidance (Section 6).

Our contributions are:

- **dfCAR metric:** displacement-filtered CAR that separates genuine avoidance from action collapse, robust across $\tau_d \in [2, 200]$ mm.
- **Action-collapse diagnosis:** 100+ experiments showing $>99\%$ of Pi0 safe episodes are action-collapse artifacts.
- **EE-only SDF is counterproductive:** oracle refinement *reduces* CAR below unrefined baselines by redistributing collisions to unmonitored arm links.
- **Demonstration sufficiency:** collision-free demonstrations account for nearly all CAR gains (~ 26 pp); SDF shows no significant benefit across $\lambda \in [0.005, 0.1]$ ($p > 0.05$ for $\lambda \in \{0.01, 0.1\}$, 3 seeds).

2 Related Work

Vision-Language-Action Models. VLA models combine pretrained vision-language backbones with action prediction for generalist manipulation [8, 6, 9, 10]. Recent work enhances spatial reasoning through auxiliary 3D supervision [11, 12, 13], but safety considerations remain largely absent:

most VLAs train on collision-free demonstrations and evaluate task success alone, with no mechanism to distinguish genuine avoidance from incidental collision-free trajectories.

Safety in Robot Learning. Safety-aware approaches span training-time constrained RL [3] and test-time filters: control barrier functions [2], collision risk gating [5], inverse kinematics QP [14], and diffusion-process collision costs [15, 4]. SafeDiffuser [16] integrates barrier functions into diffusion denoising but does not learn intrinsic collision awareness. A recent survey [1] categorizes these into runtime shields, constrained optimization, and representation-based methods. Common to all is that safety is imposed externally. This work tests whether auxiliary SDF supervision can produce intrinsic collision awareness.

Signed Distance Functions in Robotics. SDFs encode minimum surface distance, with sign indicating free space or penetration. Neural SDFs [17] enable continuous shape representation; in robotics, SDFs serve perception [18], planning [19, 20], and differentiable collision detection [21]. CDF [22] and CSSDF-Net [23] extend SDFs to configuration space for manipulation planning. RDF [24] represents robot geometry as composable per-link SDFs, enabling whole-body collision queries—directly relevant to the full-body gap our work identifies. Neural SDF-based MPC [25] is effective in classical pipelines, but integration into end-to-end learned policies remains unexplored.

SDF integration pathways. SDF-based geometric supervision can enter VLA policies through three mechanisms: (a) *training-time auxiliary loss* on the EE trajectory (tested here); (b) *constrained optimization*, where SafeVLA [3] embeds safety constraints via Lagrangian optimization; and (c) *runtime filtering*, where VLSA [2] applies control barrier functions at inference. Our negative finding is specific to pathway (a) and does not imply limitations for (b) or (c).

Evaluation Metrics for Safe Manipulation. SafeLIBERO [2] extends LIBERO [26] with obstacle-augmented environments and a CAR metric. However, CAR treats all collision-free episodes identically: an action-collapsed robot scores near-perfect CAR without any avoidance behavior, echoing concerns about metric reliability [7, 27] and binary success metrics under limited samples [28]. Action collapse in diffusion policies [29] systematically inflates CAR but has not been examined as a safety evaluation confounder. Concurrent work expands the landscape: HazardArena [30] evaluates semantic safety, and VLA-Risk [31] benchmarks physical robustness. We propose dfCAR (§3.2) to detect this contamination.

3 Experimental Setup and Metrics

3.1 Benchmark and Models

We evaluate on SafeLIBERO [2], the *object-pudding* suite (Level II): four tabletop pick-and-place tasks with collision hazards and per-timestep collision detection. Two VLA families are tested: **Pi0** (~2.5B) [6] with flow matching [32], and **Pi0.5** with broader pre-training. Table 1 lists all configurations. Each uses 50 episodes per task (200 total); Pi0.5 multi-seed runs use 3 seeds (600 episodes). All metrics are computed using a fixed evaluation seed (seed = 7); the reported $\pm$ values reflect variance across training seeds, not evaluation stochasticity. We report Collision Avoidance Rate (CAR), the fraction of collision-free episodes, and Task Success Rate (TSR), the fraction of episodes that complete the task [2, 3].

3.2 Displacement-Filtered CAR

Standard CAR conflates genuine avoidance with action collapse. We define an episode as *action-collapsed* when $d_{\max}(e) < \tau_d$ ($\tau_d = 0.1$ mm) and exclude such episodes:

$$\text{dfCAR} = \frac{|\{e : \text{safe}(e) \wedge d_{\max}(e) \geq \tau_d\}|}{|\{e : d_{\max}(e) \geq \tau_d\}|} \quad (1)$$

Table 1: Training configurations. Trainable parameters exclude the frozen backbone. Pi0+A1 LoRA weights were not loaded due to a parameter-name mismatch; only auxiliary heads were trained.

Configuration	Description	Trainable Params
Pi0 Full-FT	Full fine-tune, flow-matching loss	$\sim 2.5\text{B}$
Pi0 LoRA	LoRA adapters only	$\sim 33.5\text{M}$
Pi0+A1 (SDF aux)	LoRA + SDF decoder + EE extractor	$\sim 658\text{K}$
Pi0+A3 (learned SDF)	LoRA + SDF decoder w/ gradient refinement	$\sim 34.2\text{M}$
Pi0.5 (pre-trained)	No fine-tuning	0
Pi0.5 LoRA	LoRA fine-tune on Pi0.5	$\sim 33.5\text{M}$
Pi0.5+A1 (SDF aux)	LoRA + SDF aux on Pi0.5, $\lambda \in \{0.005, 0.01, 0.1\}$	$\sim 34.2\text{M}$

The displacement distribution is sharply bimodal, with collapsed episodes below $1\mu\text{m}$ and active episodes above 200mm ; any $\tau_d \in [2, 200]\text{mm}$ yields identical results (Appendix B). For Pi0 under A1, d_{max} uses obstacle displacement (A1’s parameter-loading failure produces an inconsistent EE reference frame); Pi0.5 uses standard EE displacement. The five-order-of-magnitude bimodal gap renders these measures functionally equivalent for collapse detection; alternative metrics (action variance, mean velocity, joint-space distance) would yield equivalent classification given this bimodal separation.

3.3 SDF Auxiliary Training and Oracle Refinement

We augment LoRA-adapted [33] VLAs with an SDF decoder and an EE position extractor (Appendix L, Table 15). The total loss is:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{action}} + \lambda (\mathcal{L}_{\text{SDF}} + \mathcal{L}_{\text{EE}}) \quad (2)$$

where $\lambda \in \{0.005, 0.01, 0.1\}$. In Pi0+A1, a parameter-name mismatch prevented LoRA loading and caused action collapse (TSR=0%; Table 1). To test the upper bound of EE-only refinement (Section 5), we also replace the learned SDF with ground-truth distance fields from the simulation geometry in three oracle variants (Appendix I).

4 The Action-Collapsed Safety Problem

Standard CAR assumes that collision-free episodes reflect deliberate avoidance. This assumption is systematically violated: in every Pi0 configuration tested, over 99% of collision-free episodes are action-collapse artifacts (Table 2).¹ Applying dfCAR collapses all three to $\approx 0\%$. Pi0+A1 achieves the *highest* raw CAR (91.5%) because more complete action collapse (from a parameter-name mismatch preventing LoRA loading; Table 1) yields greater CAR inflation. Pi0.5 validates the metric from the opposite direction: every safe episode exceeds 200mm displacement, and dfCAR equals raw CAR (28.5%). The displacement distribution is sharply bimodal (Appendix A), with a five-order-of-magnitude gap that makes τ_d robust.

Figure 2 reveals spatial structure in the TSR-CAR plane. Pi0-family models cluster in the upper-left (high CAR, near-zero TSR), Pi0.5 in the lower-right (high TSR, low CAR). The upper-right quadrant (high TSR and CAR) represents genuinely safe behavior; the upper-left is accessible through action collapse alone. Under dfCAR, Pi0 configurations drop to the origin: the apparent safety-capability tradeoff was an artifact of action-collapse contamination.

5 End-Effector-Only SDF Cannot Ensure Full-Body Safety

¹Pi0-family values report obstacle displacement from the original protocol; Pi0.5 reports EE displacement from the rotation-corrected protocol (§3.2). Sub-millimeter displacement indicates action collapse in both cases.

Table 2: Action-collapse decomposition on SafeLIBERO. dfCAR excludes episodes with displacement $< \tau_d = 0.1$ mm. Pi0+A1 dfCAR ≈ 0 : only 3/732 safe episodes exceed the threshold.

Method	Ep.	CAR (%)	TSR (%)	Collapsed (%)	dfCAR (%)
Pi0 Full-FT	200	83.0	0.5	99.4	≈ 0
Pi0 LoRA	200	82.0	0.0	~ 99	≈ 0
Pi0+A1 (SDF)	800	91.5	0.0	99.6	≈ 0
Pi0.5	200	28.5	71.5	0.0	28.5

Table 3: Oracle SDF refinement with ground-truth distance fields. All oracle variants achieve lower CAR than unrefined baselines. Per-task breakdown in Appendix F.

Method	CAR (%)	Collapsed (%)	Refines/ep
LoRA baseline (no refine)	82.0	99.4	—
Full-FT baseline (no refine)	83.0	—	—
Oracle-adaptive (GT, dyn. α)	75.0	98.3	14.2
Oracle-fixed (GT, $\alpha=0.005$)	73.1	98.3	14.8
Oracle-feature (GT, feat. space)	76.9	96.7	241.7
A3 learned SDF	77.1	—	0.0

136 Could a perfect SDF signal enable effective EE-
137 only refinement? We replace the learned SDF
138 with ground-truth distance fields from the sim-
139 ulation geometry (§3.3). Table 3 and Figure 3
140 show that all three oracle variants achieve *lower*
141 CAR than unrefined baselines: the best oracle
142 (76.9%) falls 5.1 points below LoRA (82.0%).
143 Refinement with perfect distance information
144 produces more collisions, not fewer.

145 Per-task analysis reveals the mechanism (Ap-
146 pendix F). On the one task where the EE is the
147 primary collision body, oracle refinement helps
148 (+9 points). On the remaining three, the SDF
149 gradient pushes the EE away from obstacles but
150 rotates intermediate arm links into obstacle vol-
151 umes; the net effect is one collision traded for
152 another. The A3 learned SDF never triggers re-
153 finement (0 steps per episode). EE-only refine-
154 ment is structurally counterproductive: joint-
155 space coupling between EE displacement and
156 arm-link rotation redistributes collisions to unmonitored links regardless of signal quality. Both
157 tested baselines are largely action-collapsed (Pi0+A1 TSR $\approx 0\%$, LoRA 99.4% collapsed); the struc-
158 tural finding should hold regardless of base-policy activity level, but validation on a functional policy
159 remains future work.

160 6 Avoidance Source Diagnosis

161 Fine-tuning on collision-free demonstrations produces genuine collision avoidance: LoRA-adapted
162 Pi0.5 raises CAR by ~ 26 percentage points with zero action-collapsed episodes. Does SDF geomet-
163 ric supervision contribute additional avoidance beyond what demonstrations provide? We compare




Figure 2: TSR vs. CAR for all configurations. Pi0-family models cluster in the upper-left (high CAR from inaction), Pi0.5 in the lower-right (high TSR with collision exposure), fine-tuned Pi0.5 in the center. The upper-right quadrant (high TSR and CAR) represents genuinely safe behavior. Error bars: ± 1 std across seeds.

Table 4: Avoidance source attribution on SafeLIBERO. SDF shows no statistically significant CAR improvement beyond LoRA-only at the tested scale ($p > 0.05$ for $\lambda \in \{0.01, 0.1\}$). All fine-tuned models exhibit 0% action-collapse rate and displacement > 0.37 m.

Method	Seeds	TSR (%)	CAR (%)	EE Disp (m)
Pi0.5 baseline	1	71.5	28.5	0.367
LoRA-only (no SDF)	3	51.5 ± 0.0	54.67 ± 2.02	0.391
LoRA+SDF $\lambda=0.005$	3	54.0 ± 3.0	54.5 ± 2.29	0.387
LoRA+SDF $\lambda=0.01$	3	53.8 ± 1.5	56.33 ± 1.61	0.389
LoRA+SDF $\lambda=0.1$	3	50.2 ± 3.9	57.00 ± 2.18	0.394

LoRA-only against LoRA+SDF across a $20\times$ range of loss weights ($\lambda \in \{0.005, 0.01, 0.1\}$) on Pi0.5, all evaluated at the 30K-step checkpoint.

Table 4 shows that collision-free demonstrations alone raise CAR by ~ 26 percentage points ($28.5\% \rightarrow 54.67\%$) with zero action-collapsed episodes and displacement above 0.37 m, confirming genuine avoidance. Adding SDF supervision does not yield statistically significant improvement at the tested scale: LoRA-only CAR (54.67%) is statistically indistinguishable from all SDF conditions including $\lambda = 0.1$ (57.00% , $p = 0.131$; Appendix G). A Phase 2 ablation ($\lambda = 0$ after step 5K) yields $\text{CAR} = 57.5\%^2$, suggesting that SDF gradients do not produce a lasting behavioral effect. The per-task CAR ordering is preserved across all conditions (chocolate-pudding highest, bbq-sauce lowest) and invariant to λ across the tested range, indicating that avoidance patterns reflect the demonstration distribution rather than geometric supervision.




Figure 3: CAR for oracle SDF refinement versus unrefined baselines. All GT oracle variants underperform baselines despite perfect distance information.

CAR-TSR trade-off. Fine-tuning degrades TSR by ~ 18 – 20 pp ($71.5\% \rightarrow 51.5$ – 54.0%), a consequence of distributional shift from the pre-trained prior rather than learned paralysis (all models retain > 0.37 m displacement and 0% action-collapsed episodes). The SSR/TSR ratio increases from 0.39 (baseline) to 0.51–0.60 (Appendix E): among task-completing episodes, a larger fraction avoids collisions after fine-tuning.

Dose-response at higher loss weights. Increasing λ by $20\times$ to 0.1 yields $\text{CAR} = 57.00\% \pm 2.18$ (3 seeds: 58.0%, 58.5%, 54.5%), within the null distribution ($p = 0.131^3$; Figure 4). The binding constraint is not signal magnitude but the structural mismatch between EE monitoring and full-body collision surfaces.

Robustness to null construction. Pooling LoRA-only and $\lambda = 0.005$ seeds as null (Null B, $\binom{9}{3} = 84$ permutations, $p = 0.131$) or using only LoRA-only seeds (Null A, $\binom{6}{3} = 20$ permutations, $p = 0.300$) both yield non-significance at $\alpha = 0.05$.

Statistical power. With $n = 3$ seeds per group and $\sigma \approx 2.0$ pp (LoRA cross-seed standard deviation), a two-sample t -test at $\alpha = 0.05$ (one-tailed) achieves




Figure 4: Effect of SDF loss weight λ on CAR and TSR. $\lambda = 0$: LoRA-only. The $\lambda \in \{0.01, 0.1\}$ values remain within the null band ($p > 0.05$), consistent with the absence of statistically significant improvement across a $20\times$ loss-weight range.

²At the 18K checkpoint (13K steps of Phase 2 training, single seed s2); full 30K training did not complete due to disk constraints.

³One-tailed $p = 0.071$; even with the directional prior that SDF should not *decrease* CAR, the result remains non-significant.

80% power only for effects ≥ 5.0 pp. The observed effect of ~ 2.3 pp falls well below this minimum detectable effect (MDE), so the study is underpowered to detect a benefit of this magnitude.

Evaluation-seed consistency. Under matched evaluation conditions (Appendix K), SDF $\lambda = 0.1$ yields higher CAR in all three paired eval seeds (+1.5, +2.5, +5.5 pp), with lower cross-seed variance ($\sigma = 0.9$ pp vs. 2.0 pp for LoRA). The observed effect (~ 2.3 pp) falls below the MDE (~ 5 pp); whether this directional pattern reflects a genuine small benefit or residual confounding cannot be resolved at $n = 3$.

7 Discussion and Conclusion

Across 100+ experiments and 6000+ evaluation episodes on two VLA model families, three findings converge: CAR-based VLA safety evaluation on SafeLIBERO is structurally unreliable. Action collapse renders reported CAR illusory (§3.2); EE-only oracle refinement is counterproductive (§5); and collision-free demonstrations account for nearly all measured avoidance gains, with EE-only SDF contributing no statistically significant improvement (§6).

Why EE-only SDF fails. Two independent causes explain the null result. (i) *Demonstration data subsumes the SDF signal.* Collision-free demonstrations implicitly encode spatial avoidance; the ~ 26 pp CAR gain arises from this distributional shift, with per-task ordering invariant across all SDF conditions (§6). Even at $\lambda = 0.1$, CAR remains within the null distribution ($p = 0.131$). (ii) *EE monitoring cannot prevent full-body collisions.* The SDF gradient acts on a single point while collisions occur along intermediate links; ground-truth oracle refinement *reduces* CAR by 5–10 pp (§5) because pushing the EE away rotates unmonitored links through obstacles. These causes are independent; effective geometric safety requires both full-body supervision and training signals beyond demonstrations. Constrained-optimization approaches (e.g., SafeVLA) address this through a fundamentally different paradigm—directly constraining violation probability—while runtime filtering sidesteps training altogether.

Implications for evaluation. dfCAR corrects action-collapse contamination using only EE position logging, already available in most evaluation pipelines. The bimodal displacement distribution makes τ_d robust across five orders of magnitude. We recommend reporting both CAR and dfCAR; the gap quantifies contamination across a $\sim 4000\times$ parameter range (Pi0 Full-FT ~ 2.5 B to Pi0+A1 ~ 658 K trainable). More broadly, safety may need to operate at the planning or joint-space constraint layer, since any EE-centric loss inherits the same blind spot.

Limitations and scope. All experiments use MuJoCo simulation (SafeLIBERO, two VLA families, four pick-and-place tasks). The action-collapse mechanism is structurally general—any policy prone to action collapse will inflate CAR—but its prevalence across other benchmarks and real hardware remains untested. The EE-only SDF null result is specific to auxiliary loss integration and should not be extrapolated to alternative pathways (Section 2). The sim-to-real gap introduces uncertainty: collision detection precision differs on physical hardware and sensor noise may affect dfCAR differently. Statistical power is limited by $n=3$ seeds per condition (MDE ≈ 5 pp at 80% power). Oracle experiments used largely action-collapsed models; validation on actively moving policies remains future work. Full-body SDF remains untested.

Future work. Three directions follow: (1) full-body SDF supervision along the kinematic chain; (2) cross-architecture evaluation on OpenVLA, RT-2, and other VLA families; (3) planning-layer safety constraints in joint space rather than EE space.

Broader impact. dfCAR is a diagnostic complement to standard CAR, not a replacement; practitioners should report both to distinguish genuine avoidance from action-collapse artifacts. Our

finding that EE-only SDF is redundant at the tested scale is a methodology insight, not a verdict against safety-aware training—geometric supervision may prove effective through alternative pathways or full-body collision surfaces.

References

- [1] Q. Li, B. Yin, W. Huang, R. Liu, B. Zou, R. Yu, J. Ye, W. Yu, and X. Wang. Vision-language-action safety: Threats, challenges, evaluations, and mechanisms. *arXiv preprint arXiv:2604.23775*, 2026.
- [2] S. Hu, Z. Liu, S. Liu, J. Cen, Z. Meng, and X. He. VLSA: Vision-language-action models with plug-and-play safety constraint layer. *arXiv preprint arXiv:2512.11891*, 2025.
- [3] B. Zhang, Y. Zhang, J. Ji, Y. Lei, J. Dai, Y. Chen, and Y. Yang. SafeVLA: Towards safety alignment of vision-language-action model via constrained learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025.
- [4] R. H. Yang, X. Zhao, L. M. Brunswic, M. Alban, M. Clemente, T. Cao, J. Jin, and A. Rasouli. CAPE: Context-aware diffusion policy via proximal mode expansion for collision avoidance. *arXiv preprint arXiv:2511.22773*, 2025.
- [5] X. Zhai, B. Ou, Q. Yu, C. Hao, and Y. Liu. CoFreeVLA: Collision-free dual-arm manipulation via vision-language-action model and risk estimation. *arXiv preprint arXiv:2601.21712*, 2026.
- [6] K. Black, N. Brown, D. Driess, A. Esmail, M. Equi, C. Finn, N. Fusai, L. Groom, K. Hausman, B. Ichter, et al. π_0 : A vision-language-action flow model for general robot control. *arXiv preprint arXiv:2410.24164*, 2024.
- [7] J. Karwowski, O. Hayman, X. Bai, K. Kiendlhofer, C. Griffin, and J. Sherburn. Goodhart’s law in reinforcement learning. In *International Conference on Learning Representations (ICLR)*, 2024.
- [8] A. Brohan, N. Brown, J. Carbajal, Y. Chebotar, X. Chen, K. Choromanski, T. Ding, D. Driess, A. Dubey, C. Finn, et al. RT-2: Vision-language-action models transfer web knowledge to robotic control. In *Conference on Robot Learning (CoRL)*, 2023.
- [9] C. Chi, S. Feng, Y. Du, Z. Xu, E. Cousineau, B. Burchfiel, and S. Song. Diffusion policy: Visuomotor policy learning via action diffusion. In *Robotics: Science and Systems (RSS)*, 2023.
- [10] M. J. Kim, K. Pertsch, S. Karamcheti, T. Xiao, A. Balakrishna, S. Nair, R. Rafailov, E. Foster, G. Lam, M. Nasiriany, et al. OpenVLA: An open-source vision-language-action model. In *Conference on Robot Learning (CoRL)*, 2024.
- [11] F. Li, W. Song, H. Zhao, J. Wang, P. Ding, D. Wang, L. Zeng, and H. Li. Spatial forcing: Implicit spatial representation alignment for vision-language-action model. *arXiv preprint arXiv:2510.12276*, 2025.
- [12] Y. Li, Y. Chen, M. Zhou, H. Li, Z. Zhang, and D. Zhao. QDepth-VLA: Quantized depth prediction as auxiliary supervision for vision-language-action models. *arXiv preprint arXiv:2510.14836*, 2025.
- [13] T. Yuan, Y. Liu, C. Lu, Z. Chen, T. Jiang, and H. Zhao. DepthVLA: Enhancing vision-language-action models with depth-aware spatial reasoning. *arXiv preprint arXiv:2510.13375*, 2025.
- [14] A. Spalter, M. Roberts, and L. M. Hiatt. Towards online safety corrections for robotic manipulation policies. *arXiv preprint arXiv:2409.08233*, 2024.

- [15] H. Li, Q. Feng, Z. Zheng, J. Feng, Z. Chen, and A. Knoll. Language-guided object-centric diffusion policy for collision-aware robotic manipulation. In *IEEE International Conference on Robotics and Automation (ICRA)*, 2025.
- [16] W. Xiao, T.-H. Wang, C. Gan, R. Hasani, M. Lechner, and D. Rus. SafeDiffuser: Safe planning with diffusion probabilistic models. In *International Conference on Learning Representations (ICLR)*, 2025.
- [17] J. J. Park, P. Florence, J. Straub, R. Newcombe, and S. Lovegrove. DeepSDF: Learning continuous signed distance functions for shape representation. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019.
- [18] J. Ortiz, A. Clegg, J. Dong, E. Sucar, D. Novotny, M. Zollhöfer, and M. Mukadam. iSDF: Real-time neural signed distance fields for robot perception. In *Robotics: Science and Systems (RSS)*, 2022.
- [19] W. Huang, C. Wang, R. Zhang, Y. Li, J. Wu, and L. Fei-Fei. VoxPoser: Composable 3D value maps for robotic manipulation with language models. In *Conference on Robot Learning (CoRL)*, 2023.
- [20] V. Vasilopoulos, S. Garg, P. Piacenza, J. Huh, and V. Isler. RAMP: Hierarchical reactive motion planning for manipulation tasks using implicit signed distance functions. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2023.
- [21] Y. Zhi, N. Das, and M. Yip. DiffCo: Auto-differentiable proxy collision detection with multi-class labels for safety-aware trajectory optimization. *IEEE Transactions on Robotics*, 38(5): 2668–2685, 2022.
- [22] Y. Li, X. Chi, A. Razmjoo, and S. Calinon. Configuration space distance fields for manipulation planning. In *Robotics: Science and Systems (RSS)*, 2024.
- [23] H. Chen, Y. Zhou, Y. Zhou, and H. Wang. CSSDF-Net: Safe motion planning based on neural implicit representations of configuration space distance field. *arXiv preprint arXiv:2603.18669*, 2026.
- [24] Y. Li, Y. Zhang, A. Razmjoo, and S. Calinon. Representing robot geometry as distance fields: Applications to whole-body manipulation. In *IEEE International Conference on Robotics and Automation (ICRA)*, 2024.
- [25] M. Jacquet, M. Harms, and K. Alexis. Neural NMPC through signed distance field encoding for collision avoidance. *The International Journal of Robotics Research*, 2025.
- [26] B. Liu, Y. Zhu, C. Gao, Y. Feng, Q. Liu, Y. Zhu, and P. Stone. LIBERO: Benchmarking knowledge transfer for lifelong robot learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [27] S. Gururangan, S. Swayamdipta, O. Levy, R. Schwartz, S. Bowman, and N. A. Smith. Annotation artifacts in natural language inference data. In *Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*, 2018.
- [28] D. Snyder, A. Badithela, N. Matni, G. Pappas, A. Majumdar, M. Itkina, and H. Nishimura. Beyond binary success: Sample-efficient and statistically rigorous robot policy comparison. *arXiv preprint arXiv:2603.13616*, 2026.
- [29] Z. Dong, Y. Liu, Y. Li, H. Zhao, and J. Hao. Conditioning matters: Training diffusion policies is faster than you think. *arXiv preprint arXiv:2505.11123*, 2025.
- [30] Z. Chen, Y. Gao, L. Wang, Y. Zhao, Y. Liu, J. Li, X. Zheng, Z. Wu, C. Wang, X. Ma, and Y.-G. Jiang. HazardArena: Evaluating semantic safety in vision-language-action models. *arXiv preprint arXiv:2604.12447*, 2026.

- [31] Anonymous. VLA-Risk: Benchmarking vision-language-action models with physical robustness. In *International Conference on Learning Representations (ICLR)*, 2026. Under review.
- [32] Y. Lipman, R. T. Q. Chen, H. Ben-Hamu, M. Nickel, and M. Le. Flow matching for generative modeling. In *International Conference on Learning Representations (ICLR)*, 2023.
- [33] E. J. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, and W. Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations (ICLR)*, 2022.

A Displacement Distribution




Figure 5: Displacement distribution across evaluation episodes. Pi0-family models (blue) cluster below $1 \mu\text{m}$, while Pi0.5 episodes (orange) exceed 200 mm . The five-order-of-magnitude gap makes τ_d robust to any value in $[2, 200] \text{ mm}$.

B Displacement Threshold Sensitivity Analysis

To validate the robustness of our displacement threshold τ_d , we evaluate the action-collapse rate across four orders of magnitude. Table 5 shows that the action-collapse pattern persists across all reasonable threshold choices: A1-LoRA exhibits near-total action collapse regardless of threshold, while Pi0.5+A1 maintains 0% action-collapse rate even at the most conservative setting.

Table 5: Action-collapse rate (% of episodes with $d_{\max} < \tau_d$) across threshold values.

τ_d	A1-LoRA (Pi0)	Pi0.5 Baseline	Pi0.5+A1 (SDF aux)
0.001 mm ($1 \mu\text{m}$)	73.9	0.0	0.0
0.1 mm	99.6	0.0	0.0
1.0 mm	100.0	0.0	0.0
10 mm	100.0	0.0	0.0

The choice of $\tau_d = 0.1 \text{ mm}$ represents a conservative threshold: displacement below this value is physically negligible in manipulation tasks where typical action magnitudes are $1\text{--}10 \text{ mm}$ per step.

Even at $\tau_d = 1 \mu\text{m}$, 73.9% of A1-LoRA safe episodes are classified as action-collapsed, while all Pi0.5 episodes remain above threshold at every level tested.

Table 6 extends this analysis to dfCAR across fine-tuned Pi0.5 configurations. All eight model-seed combinations yield identical dfCAR for $\tau_d \in [2, 200]$ mm. At $\tau_d = 300$ mm, 0–2 episodes per seed fall below threshold (displacement 200–300 mm), causing ≤ 0.5 pp deviation. Our conclusions hold across the entire tested range.

Table 6: dfCAR (%) across displacement thresholds for fine-tuned Pi0.5 configurations ($n = 200$ episodes per seed). Values are identical across $\tau_d \in [2, 200]$ mm; at $\tau_d = 300$ mm, at most 2 episodes are excluded.

Condition	τ_d (mm)							
	2	5	10	20	50	100	200	300
LoRA s0	56.5	56.5	56.5	56.5	56.5	56.5	56.5	56.5
LoRA s1	52.5	52.5	52.5	52.5	52.5	52.5	52.5	52.5
SDF $\lambda=0.005$ s0	57.0	57.0	57.0	57.0	57.0	57.0	57.0	57.1
SDF $\lambda=0.005$ s1	54.0	54.0	54.0	54.0	54.0	54.0	54.0	54.0
SDF $\lambda=0.005$ s2	52.5	52.5	52.5	52.5	52.5	52.5	52.5	52.3
SDF $\lambda=0.01$ s0	57.0	57.0	57.0	57.0	57.0	57.0	57.0	57.0
SDF $\lambda=0.01$ s1	57.5	57.5	57.5	57.5	57.5	57.5	57.5	57.1
SDF $\lambda=0.01$ s2	54.5	54.5	54.5	54.5	54.5	54.5	54.5	54.0

C Per-Task TSR Breakdown

Table 7 decomposes TSR by task for the avoidance source attribution experiments (Section 6). All tasks retain meaningful success rates after fine-tuning, with Chocolate Pudding most resilient and Orange Juice most affected.

Table 7: Per-task TSR (%) at 30K training steps. Baseline reports the unified seed-7 re-evaluation (aggregate TSR = 71.5%). LoRA-only reports 3-seed means (re-evaluated with unified seed-7 protocol); SDF conditions report 3-seed means.

Task	Pi0.5 Base (seed = 7)	LoRA-only (3 seeds)	SDF $\lambda=0.005$ (3 seeds)	SDF $\lambda=0.01$ (3 seeds)
Orange Juice	74.0	30.0	33.3	24.7
Choco. Pudding	80.0	70.7	78.0	76.7
Milk	66.0	44.7	35.3	48.7
BBQ Sauce	66.0	60.7	69.3	65.3
Mean	71.5	51.5	54.0	53.8

D Per-Task CAR Breakdown

Table 8 decomposes CAR by task for all fine-tuned conditions, supporting the claim that per-task ordering is preserved across SDF loss weights (§6). Chocolate Pudding is consistently highest and BBQ Sauce consistently lowest across all conditions, including $\lambda = 0.1$. The intermediate tasks (Orange Juice and Milk) exchange rank between seeds but remain bounded by the extremes. This invariance indicates that avoidance patterns reflect the spatial statistics of collision-free demonstrations rather than geometric supervision.

For the $\lambda = 0.1$ condition, seed 1 records `avg_max_obstacle_displacement` of 0.0553 m (per-task: OJ=0.033, CP=0.055, Milk=0.041, BBQ=0.093), comparable to other fine-tuned conditions and consistent with active manipulation rather than action-collapsed behavior.

Table 8: Per-task CAR (%) at 30K steps. $\lambda=0.005$, $\lambda=0.01$, and $\lambda=0.1$ report 3-seed means alongside individual seeds. Per-task ordering (CP highest, BBQ lowest) is preserved across all conditions.

Condition	OJ	CP	Milk	BBQ	Agg.
LoRA s0	56	74	70	26	56.5
LoRA s1	54	68	72	16	52.5
LoRA s2	56	76	62	26	55.0
SDF $\lambda=0.005$ (3-seed)	54.0	72.7	71.3	20.0	54.5
SDF $\lambda=0.01$ (3-seed)	57.3	72.7	72.0	23.3	56.3
SDF $\lambda=0.1$ s0	60	74	68	30	58.0
SDF $\lambda=0.1$ s1	66	72	70	26	58.5
SDF $\lambda=0.1$ s2	56	76	70	16	54.5
SDF $\lambda=0.1$ (3-seed)	60.7	74.0	69.3	24.0	57.0




Figure 6: Per-task CAR fingerprint across conditions. Each bar shows the 3-seed mean; dots indicate individual seeds. The task ordering (chocolate.pudding highest, bbq.sauce lowest) is preserved across all λ values, confirming that avoidance patterns reflect the demonstration distribution rather than geometric supervision. LoRA-only: 3-seed mean.

367 E Safe Success Rate Analysis

368 Table 9 reports the Safe Success Rate ($SSR = TSR \cap CAR$), the fraction of episodes achieving both
 369 task completion and collision avoidance. The SSR/TSR ratio quantifies the proportion of successful
 370 episodes that are also collision-free, controlling for task-completion rate.

Table 9: SSR and SSR/TSR ratio across conditions. The ratio increase from 0.39 (baseline) to 0.51–0.60 (fine-tuned) indicates that CAR gains reflect learned safer execution, not reduced task exposure.

Condition	TSR (%)	CAR (%)	SSR (%)	SSR/TSR
Pi0.5 baseline	71.5	28.5	28.0	0.39
LoRA-only (3 seeds)	51.5	54.67	27.7	0.54
SDF $\lambda=0.005$ (3 seeds)	54.0	54.5	27.8	0.51
SDF $\lambda=0.01$ (3 seeds)	53.8	56.33	30.2	0.56
SDF $\lambda=0.1$ (3 seeds)	50.2	57.00	30.2	0.60

F Per-Task Oracle Analysis

Table 10 shows per-task CAR for oracle refinement. On task 1 (`chocolate_pudding`), where the EE is the primary collision body and the LoRA baseline achieves only 46.0% CAR, Oracle-adaptive improves to 55.0% and Oracle-feature to 47.5%. On the remaining tasks, all oracle variants degrade: Oracle-adaptive reduces task 0 (`orange_juice`) from 96.0% to 82.5% and task 2 (`milk`) from 98.0% to 77.5%. The SDF gradient pushes the EE away from the nearest surface, but the resulting joint-space displacement rotates intermediate arm links into obstacle volumes.

The feature-space oracle fires on nearly every timestep (241.7 refinements/episode vs. 14.2 for Cartesian variants) but with gradient magnitudes two orders of magnitude smaller (0.004–0.006 vs. 0.18–0.49), and CAR (76.9%) shows no advantage despite $17\times$ more steps. The A3 learned SDF never triggers refinement (0 steps/episode), as its outputs avoid the activation threshold.

Table 10: Per-task CAR (%) for oracle SDF refinement variants. EE-only refinement helps only on task 1 where the EE is the primary collision body; on other tasks it trades one collision site for another.

Method	Task 0	Task 1	Task 2	Task 3	Agg.
LoRA baseline	96.0	46.0	98.0	88.0	82.0
Oracle-adaptive	82.5	55.0	77.5	85.0	75.0
Oracle-fixed	85.0	50.0	77.5	80.0	73.1
Oracle-feature	90.0	47.5	82.5	87.5	76.9

G Null Distribution for SDF Loss Weight Analysis

We construct a null distribution by pooling the three LoRA-only seeds and three $\lambda = 0.005$ seeds into six per-seed CAR measurements: 56.5% (LoRA s0, re-evaluated), 52.5% (LoRA s1), 55.0% (LoRA s2), 57.0% ($\lambda=0.005$ s0), 54.0% ($\lambda=0.005$ s1), 52.5% ($\lambda=0.005$ s2), yielding mean 54.58% and standard deviation 1.93% (Bessel-corrected, $n-1 = 5$). We treat these as exchangeable because SDF at $\lambda = 0.005$ produces no separation from LoRA-only. The $\lambda = 0.01$ mean of 56.33% falls within 0.9σ of this null ($p > 0.1$), consistent with seed-to-seed variation rather than a treatment effect. The $\lambda = 0.1$ 30K mean of 57.00% (3 seeds: 58.0%, 58.5%, 54.5%) yields a two-tailed permutation $p = 0.131$ (one-tailed $p = 0.071$). The result is non-significant at conventional thresholds ($\alpha = 0.05$).

The permutation test makes no distributional assumptions. Pooling the six null values {56.5, 52.5, 55, 57, 54, 52.5} and three $\lambda = 0.1$ values {58, 58.5, 54.5} into nine measurements, we enumerate all $\binom{9}{3} = 84$ three-element combinations and compute the fraction whose mean ≥ 57.00 . Six combinations meet this threshold, yielding an exact one-tailed permutation $p = 6/84 \approx 0.071$. For the two-tailed test, we additionally count combinations on the lower tail with equal or greater distance from the pooled mean, yielding $p = 11/84 \approx 0.131$. With the third seed ($s_2 = 54.5\%$) pulling the $\lambda = 0.1$ mean from 58.25% (2 seeds) down to 57.00%, the marginal signal observed at $n = 2$ dissolves into seed-to-seed variance.

Alternative null construction (Null A). If only the three LoRA-only seeds are used as the null pool ({56.5, 52.5, 55.0}), we obtain $\binom{6}{3} = 20$ permutations for each comparison. For $\lambda = 0.1$ (treatment mean 57.00%): one-tailed $p = 3/20 = 0.150$, two-tailed $p = 6/20 = 0.300$. For $\lambda = 0.01$ (treatment mean 56.33%): one-tailed $p = 4/20 = 0.200$, two-tailed $p = 8/20 = 0.400$. Under the broader Null B pool, $\lambda = 0.01$ yields one-tailed $p = 10/84 \approx 0.119$, two-tailed $p = 22/84 \approx 0.262$. All tests are non-significant at $\alpha = 0.05$; the narrower Null A yields *larger* p -values because the LoRA-only pool already contains a high seed (56.5%), making the treatment means less extreme by comparison.

A Phase 2 ablation reinforces this conclusion. Setting $\lambda = 0$ during Phase 2 (steps 5K–18K), so that SDF supervision is active only during auxiliary head warmup (Phase 1, steps 0–5K), yields CAR of 57.5% at the 18K checkpoint (13K steps of Phase 2 training, single seed, s2).⁴ This value lies within the null range, indicating that SDF training produces no lasting behavioral effect on collision avoidance.

H Step-wise Checkpoint Analysis for $\lambda = 0.1$

Table 11 reports per-seed CAR at seven checkpoints for $\lambda = 0.1$. Seeds 0 and 1 are available at 16K, 20K, 22K, and 30K; 24K–28K report seed-0 only; seed 2 is available at 30K only. The seed-0 curve shows no monotonic trend through 28K. At 30K, the addition of seed 2 (54.5%) pulls the three-seed mean to 57.00%, within the null band ($p = 0.131$).

Table 11: Per-seed CAR (%) at training checkpoints for $\lambda = 0.1$ (200 episodes each). 24K–28K: seed-0 only. The 30K checkpoint includes all three seeds; the three-seed mean (57.00%) remains within the null band ($p = 0.131$, two-tailed permutation).

Checkpoint	Seed 0	Seed 1	Seed 2	Mean
16K steps	55.5	56.0	–	55.75
20K steps	55.0	56.0	–	55.5
22K steps	60.5	56.5	–	58.5
24K steps	52.5	–	–	–
26K steps	53.5	–	–	–
28K steps	54.5	–	–	–
30K steps	58.0	58.5	54.5	57.00

Individual seeds exhibit checkpoint-to-checkpoint CAR fluctuations of up to 8 pp (seed 0: 52.5–60.5%), larger than the cross-seed standard deviation at any fixed checkpoint (≤ 3 pp). This is consistent with binomial sampling noise at $n = 50$ episodes per task ($\sigma \approx 7$ pp per task for $p \approx 0.5$). At 30K, the third seed ($s_2 = 54.5\%$) falls well within the null range, pulling the three-seed mean to 57.00% and rendering the earlier two-seed departure non-significant ($p = 0.131$). The seed-averaged curve shows no monotonic trend, confirming that SDF loss weight does not produce a progressive training signal.

I Oracle SDF Refinement Details

All oracle variants replace the learned SDF decoder with ground-truth distance fields computed from the simulation geometry. Each is evaluated for 160 episodes (4 tasks $\times$ 40).

Oracle-adaptive. At each timestep, if the ground-truth SDF at the end-effector position falls below a safety margin, the action is adjusted by a step in the negative SDF gradient direction. The step size α is scaled inversely with distance: smaller distances produce larger corrections. This mimics an adaptive potential field controller.

Oracle-fixed. Identical to Oracle-adaptive except the step size is fixed at $\alpha = 0.005$ regardless of distance. This isolates the effect of step-size adaptation from the refinement geometry.

Oracle-feature. Instead of modifying the action directly in Cartesian space, this variant applies the SDF gradient in the policy’s latent feature space via backpropagation. The gradient of the ground-truth SDF value with respect to the intermediate features is computed and used to perturb the features before the action head decodes them. This tests whether geometric corrections are more effective when applied in the model’s learned representation space. In practice, the feature-space variant fires

⁴Full 30K training did not complete due to disk constraints; this result uses the last complete checkpoint at step 18K.

on nearly every timestep (241.7 refinements per episode vs. 14.2 for Cartesian variants) but with gradient magnitudes two orders of magnitude smaller (0.004–0.006 vs. 0.18–0.49 in action space).

J Cross-Scene Validation: `spatial_bowl`

To test whether the absence of EE-only SDF benefit generalizes beyond the primary `object_pudding` scene, we replicate the core comparison on `spatial_bowl`, a separate Safe-LIBERO suite with a distinct obstacle layout (moka pot, wine bottle, milk carton, red coffee mug, yellow book). Both conditions use the same architecture, hyperparameters, and two-phase training schedule (30K steps), differing only in the training scene.

Table 12 reports the results. The SDF condition ($\lambda = 0.005$) yields CAR of 11.5% versus 12.0% for LoRA-only, a -0.5 pp difference well within the noise floor, indicating that SDF is essentially identical to LoRA-only on this scene. TSR is comparable across conditions (49.0% vs. 51.5%). Both evaluations use a single training seed (seed 0) and a single evaluation seed (seed 7, 200 episodes), so statistical power is limited; nevertheless, the qualitative pattern—SDF failing to improve over LoRA-only—is consistent with the `object_pudding` findings.

Table 12: Cross-scene validation on `spatial_bowl` (single training seed, 200 episodes, 30K checkpoint). The -0.5 pp CAR difference is well within noise, consistent with the absence of SDF benefit observed on `object_pudding`.

Condition	CAR (%)	TSR (%)
LoRA-only	12.0	49.0
SDF $\lambda=0.005$	11.5	51.5

The substantially lower CAR on `spatial_bowl` relative to `object_pudding` (11–12% vs. 54–57%) reflects the greater scene difficulty: `spatial_bowl` places five cluttered obstacles around the workspace, making collision-free manipulation considerably harder. This difficulty affects both conditions equally, and the SDF condition performs comparably ($\Delta\text{CAR} = -0.5$ pp, well within noise), reinforcing the conclusion that EE-only SDF auxiliary supervision does not produce a meaningful collision avoidance benefit beyond LoRA-only fine-tuning.

K Evaluation Seed Robustness

To verify that our conclusions are not artifacts of a particular evaluation seed, we re-evaluate key model configurations under multiple evaluation seeds (200 episodes each). Table 13 reports aggregate CAR and TSR across all tested seeds.

Table 13: Eval-seed robustness: aggregate CAR and TSR (%) across evaluation seeds (200 episodes each). Dashes indicate untested seeds. CAR ranges overlap across all three conditions, confirming that SDF redundancy is not an eval-seed artifact.

Condition		Evaluation seed						Mean $\pm$ Std
		0	7	42	99	123	256	
LoRA s0	CAR	53.0	56.5	56.5	57.5	53.0	—	55.3 $\pm$ 2.1
	TSR	55.5	51.5	46.0	52.0	50.5	—	51.1 $\pm$ 3.4
SDF $\lambda=0.01$ s0	CAR	52.5	57.0	—	57.0	—	—	55.5 $\pm$ 2.6
	TSR	48.5	52.5	—	46.5	—	—	49.2 $\pm$ 3.1
SDF $\lambda=0.1$ s0	CAR	—	58.0	59.0	—	58.5	57.0	58.1 $\pm$ 0.9
	TSR	—	47.0	52.5	—	53.5	48.0	50.3 $\pm$ 3.2

LoRA s0 CAR ranges from 53.0% to 57.5% across five evaluation seeds (std=2.1 pp), consistent with binomial sampling noise at $n=50$ episodes per task. SDF $\lambda=0.01$ s0 produces near-

identical mean CAR (55.5% vs. 55.3%), confirming that the null result holds across eval conditions. SDF $\lambda=0.1$ s0 shows a slightly higher mean (58.1%), but the ranges overlap with LoRA-only, and the difference (2.8 pp, $<1.3\sigma$) remains non-significant. The per-task CAR ordering (chocolate_pudding highest, bbq_sauce lowest) is preserved across all seeds and conditions, reinforcing the demonstration-driven avoidance hypothesis.

Table 14 provides per-task CAR for the five-seed LoRA evaluation, the most complete eval-seed battery.

Table 14: Per-task CAR (%) for LoRA s0 across five evaluation seeds. The per-task fingerprint (CP highest, BBQ lowest) is invariant across all seeds.

Eval seed	OJ	CP	Milk	BBQ	Agg.
0	46	70	74	22	53.0
7	56	74	70	26	56.5
42	60	70	68	28	56.5
99	58	76	66	30	57.5
123	52	74	66	20	53.0
Mean $\pm$ Std	54.4 $\pm$ 5.5	72.8 $\pm$ 2.7	68.8 $\pm$ 3.3	25.2 $\pm$ 4.1	55.3 $\pm$ 2.1

L Training Details

Architecture. We use LoRA [33] adapters injected into the PaliGemma vision encoder (rank 16, $\alpha=16$) and the action expert (rank 32, $\alpha=32$), following the configuration of OpenPi. The SDF decoder is a 4-layer MLP (hidden dimension 256) with skip connections, taking pooled suffix features (1024-d) and query points (3-d) as input. The end-effector extractor is a 2-layer MLP (1024 $\rightarrow$ 256 $\rightarrow$ 3).

Training schedule. Training proceeds in two phases:

- **Phase 1** (steps 0–5000): Backbone and LoRA parameters are frozen. Only the SDF decoder and EE extractor are trained, allowing them to learn meaningful representations before the policy adapts.
- **Phase 2** (steps 5000–30000): LoRA parameters are unfrozen with a 100-step gradient warmup ramp. All auxiliary heads continue training jointly with the policy.

Table 15: Training hyperparameters.

Parameter	Value
Base model	Pi0.5 (OpenPi, ~ 3.5 B parameters)
Optimizer	AdamW
Learning rate	2.5×10^{-5} (cosine decay to $\text{lr}/10$)
Warmup steps	1000
Batch size	32
Total steps	30000
λ_{sdf}	{0.005, 0.01, 0.1}
λ_{ee}	0.01
SDF query points per step	64
SDF clamp δ	0.05 m
Eikonal weight	0.1
LoRA rank (vision)	16
LoRA rank (action)	32
Phase 1 duration	5000 steps

Hyperparameters.

485 **Evaluation protocol.** Each episode runs for at most 300 timesteps. The policy generates a 50-
486 step action trajectory via flow matching and executes the first 5 steps before replanning. An episode
487 terminates early upon task completion. Collision detection uses MuJoCo’s built-in contact detection:
488 any contact between the robot and a non-target object registers as a collision. For Pi0.5 evaluations,
489 camera images are rotated 180° to match the pre-training viewpoint convention; this correction was
490 verified to restore TSR from $\sim 5\%$ to $\sim 52\%$ while leaving CAR comparable. Oracle variants use
491 40 episodes per task (160 total) with a single seed; all other configurations use 50 episodes per task
492 (200 total).

493 **Data.** We use the SafeLIBERO object_pudding suite with pre-computed SDF ground truth. SDF
494 values are computed offline from the obstacle mesh using a dense 64^3 voxel grid. During training,
495 64 query points are randomly sampled per step from this grid.

496 **Computational resources.** All training and evaluation experiments are conducted on NVIDIA
497 RTX PRO 6000 GPUs (96 GB). Phase 1 (5K steps) takes approximately 6 hours on a single GPU;
498 Phase 2 (25K steps) takes approximately 32 hours. Evaluation of 200 episodes requires approxi-
499 mately 4 hours. Across 100+ experiments (training runs, evaluations, oracle variants, and ablations),
500 the total computational budget is approximately 700 GPU-hours.